

6th Seminar

IPCs

There are several types of Inter Process Communication mechanisms:


- Message queues
- Shared memory
- semaphores

IPC structures are identified by an integer identifier (non negative)

IPCs are created with the use of a key of `key_t` type and OS transforms the key into IPC identifier

Shared Memory

- A shared memory region is a portion of physical memory that is shared by multiple processes. In this region, structures can be set up by processes and others may read/write on them. Synchronization, if required, is achieved with the help of synchronization mechanisms (semaphores for example).



Shared Memory

- create a shared memory segment

```
int shmget(key_t key, size_t size, int shmflg)
```

returns the identifier of the shared memory segment associated with the value of the argument `key`.

the returned size of the segment is equal to size rounded up to a multiple of PAGE SIZE.

`shmflg` helps designate the access rights for the segment.

If `shmflg` specifies both `IPC_CREAT` and `IPC_EXCL` and a shared memory segment already exists for `key`, then `shmget()` fails with `errno` set to `EEXIST`.

Shared Memory

- Attaching to a shared memory segment

`void *shmat(int shmid, const void *shmaddr, int shmflg)`
attaches the shared memory segment identified by `shmid` to the address space of the calling process.

If `shmaddr` is `NULL`, the OS chooses a suitable (unused) address at which to attach the segment (**frequent choice**). Otherwise, `shmaddr` must be a page-aligned address at which the attachment occurs.

- Detaching from a shared memory segment

`int shmdt(const void *shmaddr)`
detaches the shared memory segment located at the address specified by `shmaddr` from the address space of the calling process.

Shared Memory

`int shmctl(int shmid, int cmd, struct shmid_ds *buf)`
performs the control operation specified by `cmd` on the shared memory segment whose identifier is given in `shmid`. The `buf` argument is a pointer to a `shmid_ds` structure:

```
struct shmid_ds {
    struct ipc_perm shm_perm ; /* Ownership and permissions */
    size_t shm_segsz ; /* Size of segment ( bytes ) */
    time_t shm_atime ; /* Last attach time */
    time_t shm_dtime ; /* Last detach time */
    time_t shm_ctime ; /* Last change time */
    pid_t shm_cpid ; /* PID of creator */
    pid_t shm_lpid ; /* PID of last shmat (2) / shmdt (2) */
    shmatt_t shm_nattch ; /* No . of current attaches */
    ...
};
```

Shared Memory

- some possible values for `cmd`:
 - `IPC_STAT`: copy information from the kernel data structure associated with `shmid` into the `shmid_ds` structure pointed to by `buf`.
 - `IPC_SET`: write the value of some member of the `shmid_ds` structure pointed to by `buf` to the kernel data structure associated with this shared memory segment, updating also its `shm_ctime` member.
 - `IPC_RMID`: mark the segment to be destroyed. The segment will be destroyed after the last process detaches it (i.e., `shm_nattch` is zero).

```
// writer.c
int main()
{
    // ftok to generate unique key
    key_t key = ftok("shmfile",65);
    // shmget returns an identifier in shmid
    int shmid = shmget(key,1024,0666|IPC_CREAT);
    // shmat to attach to shared memory
    char *str = (char*) shmat(shmid,(void*)0,0);
    printf("Write Data : ");
    gets(str);
    printf("Data written in memory: %s\n",str);
    //detach from shared memory
    shmdt(str);
    return 0;
}
```



```
// reader.c
int main()
{
    // ftok to generate unique key
    key_t key = ftok("shmfile",65);
    // shmget returns an identifier in shmid
    int shmid = shmget(key,1024,0666|IPC_CREAT);
    // shmat to attach to shared memory
    char *str = (char*) shmat(shmid,(void*)0,0);
    printf("Data read from memory: %s\n",str);
    //detach from shared memory
    shmdt(str);
    // destroy the shared memory
    shmctl(shmid,IPC_RMID,NULL);
    return 0;
}
```

Semaphores

- Semaphores were proposed by Dijkstra to manage concurrent processes by using a simple integer value, which is known as a semaphore.
 - is simply a variable which is non-negative and shared between threads.
 - is used to solve the critical section problem and to achieve process synchronization in the multiprocessing environment.
- Semaphores are of two types:
 - **Binary Semaphore** – also known as mutex lock. It can have only two values – 0 and 1. Its value is initialized to 1. It is used to implement the solution of critical section problem with multiple processes/threads.
 - **Counting Semaphore** – Its value can range over an unrestricted domain. It is used to control access to a resource that has multiple instances.

Semaphores

- There are two operations which can be used to access and change the value of the semaphore variable:
 - **P** - is also called **wait**, **sleep** or **down** operation
 - **V** - is also called **signal**, **wake-up** or **up** operation
- Both operations are **atomic** and semaphore is always initialized to value 1
- A critical section is surrounded by both operations to implement process synchronization.

```
P(semaphore s){          V(semaphore s){          P(s);
    while (s==0);          s=s+1;          // critical section
    s=s-1;                  }          V(s);
}
```

Semaphores vs Mutexes

- There is a similarity between binary semaphore and mutex. But they are not the same. The purpose of mutex and semaphore are different. Due to similarity in their implementation a mutex would be referred as binary semaphore.
- A **mutex** is a **locking mechanism** used to synchronize access to a resource. Only one task (can be a thread or process based on OS abstraction) can acquire the mutex. It means there is **ownership** associated with mutex, and only the owner can release the lock (mutex).
- A **semaphore** is a **signaling mechanism**.

Semaphores

POSIX semaphores (two forms): **named** semaphores and **unnamed** semaphores.

- A **named** semaphore is identified by a name of the form /somename; a null-terminated string of up to 251 characters consisting of an **initial slash**, followed by one or more characters, none of which are slashes. Two processes can operate on the same named semaphore by passing the same name to `sem_open`.
- The `sem_open` function creates a new named semaphore or opens an existing named semaphore. After the semaphore has been opened, it can be operated on using `sem_post` and `sem_wait`. When a process has finished using the semaphore, it can use `sem_close` to close the semaphore. When all processes have finished using the semaphore, it can be removed from the system using `sem_unlink`.

Semaphores

- An **unnamed** (memory-based) semaphore does not have a name. Instead the semaphore is placed in a region of memory that is shared between multiple threads (a **thread-shared** semaphore) or processes (a **process-shared** semaphore). A thread-shared semaphore is placed in an area of memory shared between the threads of a process, for example, a global variable. A process-shared semaphore must be placed in a shared memory region (e.g., a System V shared memory segment created using `shmget`, or a POSIX shared memory object built created using `shm_open`).
- Before use, an unnamed semaphore must be initialized using `sem_init`. It can then be operated on using `sem_post` and `sem_wait`. When is no longer required, and before the memory in which it is located is deallocated, the semaphore should be destroyed using `sem_destroy`.

Semaphores

```
static int sem_init (sem_t *sem, int pshared, unsigned value)
```

Initialize an unnamed semaphore.

```
static int sem_post (sem_t *sem)
```

Unlock a semaphore.

```
static int sem_wait (sem_t *sem)
```

Lock a semaphore.

```
int sem_timedwait (sem_t *sem, const struct timespec *abstime)
```

Similar to `sem_wait' but wait only until abstime.

```
static int sem_trywait (sem_t *sem)
```

Test whether sem is posted.

```
static int sem_getvalue (sem_t *sem, int *sval)
```

Get current value of sem and store it in sval.

Semaphores

```
static sem_t * sem_open (const char *name, int oflag, ...)
```

Open a named semaphore name with open flags oflag.

```
static int sem_close (sem_t *sem)
```

Close descriptor for named semaphore sem.

```
static int sem_unlink (const char *name)
```

Remove named semaphore name.

```
static int sem_destroy (sem_t *sem)
```

destroy an unnamed semaphore


```
//POSIX Semaphore test example using shared memory
# define SEGMENTSIZe sizeof(sem_t)    // sem_t means semaphore type
# define SEGMENTPERM 0666
int main (int argc ,char ** argv )
{
    sem_t *sp;
    int retval, id, err ;
    if ( argc <= 1) {
        printf("Need shmem id. \n");
        exit (1) ;
    }
    sscanf( argv [1], "%d", &id);    // Get id from command line
    printf(" Allocated %d\n", id);
    sp = (sem_t *) shmat (id ,( void *) 0, 0);    // Attach the segment
    if (( int ) sp == -1) {
        perror (" Attachment .");
        exit (2) ;
    }
}
```

```
retval = sem_init (sp ,1 ,1) ; // Initialize the semaphore
if ( retval != 0) {
    perror(" Couldn't initialize .");
    exit (3) ;
}
retval = sem_trywait (sp);
printf("Did trywait . Returned %d >\n", retval ); getchar ();
retval = sem_trywait (sp);
printf("Did trywait . Returned %d >\n", retval ); getchar ();
retval = sem_trywait (sp);
printf("Did trywait . Returned %d >\n", retval ); getchar ();
err = shmdt (( void *) sp); // Remove segment
if (err == -1)
    perror (" Detachment .");
return 0;
}
```

Semaphores

Între A și B sunt n linii prin care trec m trenuri, $m > n$

În gara A intră simultan maximum m trenuri care vor să ajungă în gara B. De la A spre B există simultan n linii, $m > n$. Fiecare tren intră în A la un interval aleator.

Dacă are linie liberă între A și B, o ocupă și pleacă către B, durata de timp a trecerii este una aleatoare. Să se simuleze aceste treceri. Soluțiile, una folosind variabile condiționale, cealaltă folosind semafoare, sunt prezentate în tabelul următor.

```

#define N 5
#define M 13
#define SLEEP 4

pthread_mutex_t mutcond;
pthread_cond_t cond;
int linie[N], tren[M], inA[M+1],
dinB[M+1], liniilibere;
pthread_t tid[M];
time_t start;
void t2s(int *t, int l, char *r) {
    int i;
    char n[10];
    sprintf(r, "[");
    for ( i = 0; i < l; i++) {
        sprintf(n,"%d, ",t[i]);
        strcat(r, n);}
    i = strlen(r) - 1;
    if(r[i] == ' ') r[i - 1] = 0;
    strcat(r, "]);
}

```

```

#define N 5
#define M 13
#define SLEEP 4

sem_t sem; // Asteapta / semnaleaza
eliberarea uneia din cele N linii
sem_t sem, mut; // Asigura acces
exclusiv la tabelele globale
int linie[N], tren[M], inA[M+1],
dinB[M+1];
pthread_t tid[M];
time_t start;
void t2s(int *t, int l, char *r) {
    int i;
    char n[10];
    sprintf(r, "[");
    for( i = 0; i < l; i++) {
        sprintf(n,"%d, ",t[i]);
        strcat(r, n);
    }
}

```

```

void printT(char *s, int t) {
    int i;
    char a[200],l[200],b[200];
    for(i=0; inA[i] != -1; i++);
    t2s(inA, i, a);
    t2s(linie, N, l);
    for(i=0; dinB[i] != -1; i++);
    t2s(dinB, i, b);
    printf("%s
%d\tA:%s\tLines:%s\tB:%s\ttime:
%ld\n",s,t,a,l,b,time(NULL)-start);
}

```

//rutina unui thread

```

void* trece(void* tren) {
    int i, t, l;
    t = *(int*)tren;
    sleep(1 + rand() % SLEEP); //

```

Modificati timpii de stationare

```

    i = strlen(r) - 1;
    if( r[i] == ' ') r[i - 1] = 0;
    strcat(r, "]\n");
}
void printT(char *s, int t) {
    int i;
    char a[200],l[200],b[200];
    for(i = 0; inA[i] != -1; i++);
    t2s(inA, i, a);
    t2s(linie, N, l);
    for(i = 0; dinB[i] != -1;
i++);

    t2s(dinB, i, b);
    printf("%s
%d\tA:%s\tLines:%s\tB:%s\ttime:
%ld\n",s,t,a,l,b,time(NULL)-start);
}

```

```

pthread_mutex_lock(&mutcond);
for(i=0;inA[i]!=-1;i++);
inA[i]=t;
printT("EnterA", t);
for(;liniilibere==0;)
pthread_cond_wait(&cond,&mutcond);
for(l=0;linie[l]!=-1;l++);
linie[l]=t;
liniilibere--;
for(i=0;inA[i]!=t;i++);
for(;i<M;inA[i]=inA[i+1],i++);
printT(" A ==> B", t);
pthread_mutex_unlock(&mutcond);

sleep(1 + rand() % SLEEP);

```

```

//rutina unui thread
void* trece(void* tren) {
    int i, t, l;
    t = *(int*)tren;
    sleep(1+rand()%SLEEP); //Inainte de
==> A
    sem_wait(&mut);
    for( i = 0; inA[i] != -1; i++);
    inA[i] = t;
    printT("EnterA", t);
    sem_post(&mut);
    sem_wait(&sem); // In A ocupa linia
    sem_wait(&mut);
    for(l = 0; linie[l] != -1; l++);
    linie[l] = t;

```

```

pthread_mutex_lock(&mutcond);
linie[1] = -1;
liniilibere++;
for(i=0; dinB[i] != -1; i++);
dinB[i] = t;
prinT("  OutB", t);
pthread_cond_signal(&cond);
pthread_mutex_unlock(&mutcond);

```

```

}

```

```

//main

```

```

int main(int argc, char* argv[]) {
    int i;
    start = time(NULL);
    pthread_mutex_init(&mutcond, NULL
);

    pthread_cond_init(&cond, NULL);
    liniilibere = N;

```

```

        for( i = 0; inA[i] != t; i++);
        for( ; i < M; inA[i] = inA[i +
1], i++);
        prinT(" A ==> B", t);
        sem_post(&mut);

```

```

        sleep(1+rand()%SLEEP); //

```

```

Trece trenul  A ==> B

```

```

        sem_wait(&mut);
        linie[1] = -1;
        for(i=0; dinB[i] != -1; i++);
        dinB[i] = t;
        prinT("  OutB", t);
        sem_post(&mut);

```

```

        sem_post(&sem); //In B

```

```

elibereaza linia

```

```

}

```

```

        for(i=0;i<N;linie[i]=-1,i++);
        for(i=0;i<M;tren[i]=i,i++);
        for (i=0;i<M+1;inA[i]=-
1,dinB[i]= -1, i++);

        // ce credeti despre ultimul
parametru &i?
        for(i=0;i<M;i++)
pthread_create(&tid[i],NULL,trece,
&tren[i]);
        for(i=0;i<M;i++)
pthread_join(tid[i], NULL);

        pthread_mutex_destroy(&mutcond
);
        pthread_cond_destroy(&cond);
        return 0;
}

```

```

int main(int argc, char* argv[]) {
    int i;
    start = time(NULL);
    sem_init(&sem, 0, N);
    sem_init(&mut, 0, 1);
    for(i=0;i<N;linie[i]=-1,i++);
    for(i=0;i<M;tren[i]=i,i++);
    for(i=0;i<M+1;inA[i]=-
1,dinB[i] = -1, i++);
    // ce credeti despre ultimul
parametru &i in loc de &tren[i]?
    for(i=0; i < M; i++)
pthread_create(&tid[i],NULL,trece,
&tren[i]);
    for(i=0;i<M;i++)
pthread_join(tid[i], NULL);
    sem_destroy(&sem);
    sem_destroy(&mut);
    return 0;
}

```


Barriers

- When multiple threads are working together, it might be required that threads wait for each other at a certain event or point in the program before proceeding ahead.
- Such operations can be implemented by adding a barrier in the thread. A barrier is a point where the thread is going to wait for other threads and will proceed further only when predefined number of threads reach the same barrier.
- A barrier is a variable of the type `pthread_barrier_t`.

```
int pthread_barrier_init(pthread_barrier_t *restrict  
barrier, const pthread_barrierattr_t *restrict attr, unsigned  
count);
```

- `barrier`: The variable used for the barrier
- `attr`: Attributes for the barrier, which can be set to NULL to use default attributes
- `count`: Number of threads which must wait call `pthread_barrier_wait` on this barrier before the threads can proceed further.

Barriers

- Once a barrier is initialized, a thread can be made to wait on the barrier for other threads using the function

```
int pthread_barrier_wait(pthread_barrier_t *barrier);
```

where the barrier is the same variable initialized using `pthread_barrier_init`.

- A thread will keep waiting till the "count" number of threads passed during init do not call the `wait` function on the same barrier.
- Once `pthread_barrier_wait` has been called by "count" threads, the constant `PTHREAD_BARRIER_SERIAL_THREAD` is returned to one unspecified thread and 0 is returned to each of the remaining threads. The barrier is then reset to the state it had as a result of the most recent `pthread_barrier_init()` function that referenced it..

Barriers

- A barrier can be destroyed using the function

```
int pthread_barrier_destroy(pthread_barrier_t *barrier);
```

The barrier should be destroyed only when no thread is executing a wait on the barrier

```
pthread_barrier_t  barrier; // the barrier synchronization object
void * thread1 (void *not_used){
    time_t  now;
    char    buf [27];
    time (&now);
    printf ("thread1 starting at %s", ctime_r (&now, buf));
    // do the computation
    // let's just do a sleep here...
    sleep (20);
    pthread_barrier_wait (&barrier);
    // after this point, all three threads have completed.
    time (&now);
    printf ("barrier in thread1() done at %s", ctime_r (&now,
buf));
}
```

```
void * thread2 (void *not_used) {
    time_t  now;
    char    buf [27];

    time (&now);
    printf ("thread2 starting at %s", ctime_r (&now, buf));

    // do the computation
    // let's just do a sleep here...
    sleep (40);
    pthread_barrier_wait (&barrier);
    // after this point, all three threads have completed.
    time (&now);
    printf ("barrier in thread2() done at %s", ctime_r (&now,
buf));
}
```

```
int main () {
    time_t  now;
    char buf [27];
    // create a barrier object with a count of 3
    pthread_barrier_init (&barrier, NULL, 3);
    // start up two threads, thread1 and thread2
    pthread_create(NULL, NULL, thread1, NULL);
    pthread_create(NULL, NULL, thread2, NULL);
    // at this point, thread1 and thread2 are running
    // now wait for completion
    time(&now);
    printf("main waiting for barrier at %s", ctime_r (&now, buf));
    pthread_barrier_wait(&barrier);
    // after this point, all three threads have completed.
    time(&now);
    printf("barrier in main () done at %s", ctime_r (&now, buf));
}
```